ECE-471 Selected Topics in Machine Learning
Prof. Curro
Midterm Project

Evan Bubniak, Jinhan Zhang

October 29, 2019

1   Summary

The purpose of this project is to reproduce a subset of experimental results from the paper *Understanding deep learning requires rethinking generalization* (Zhang et al. 2016), specifically Figure 1, which plots the average loss against the number of training steps for various types of example/label corruption. The experiment demonstrates that various kinds of label corruption do not prevent complete memorization of the training dataset, as long as the model is sufficiently large, and that the convergence of the model to 100% accuracy is shifted by only a constant factor when data corruption is introduced.

Our results show strong agreement with Figure 1 for the first model specified, the minified InceptionV3, and strong divergence in all other models. The two models are displayed side-by-side. Crucially, we verify the key finding of the paper, which is the constant-factor difference in convergence time when introducing data and label corruption.

Figure 1: A screenshot of the original figure as it appears in C. Zhang et al.
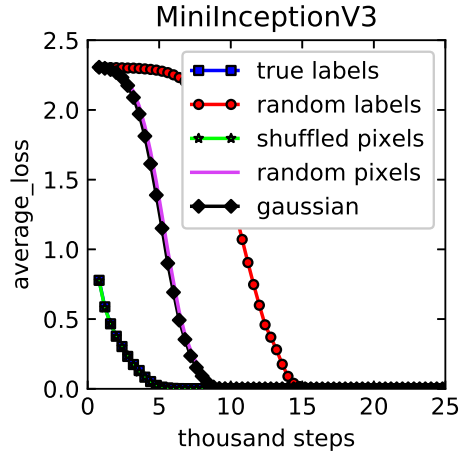

Figure 2: A plot of the average losses for the minified InceptionV3 on the CIFAR-10 dataset, subject to various types of data corruption.

2   Experimental Setup

We conducted many tests to reconstruct and verify the experimental setup. Specifically, we rebuilt each of the five models described in Table 1 of the original paper: two minified InceptionV3 models respectively with and without batch normalization,

a minified AlexNet, and two multi-layer perceptions with 512 units and, respectively, one and three hideen layers. These models were built in Keras and the Python code defining them is provided in the appendix.

The experimental setup involved, in total, the training and analysis of 25 models, corresponding to five models each trained on five datasets. Training was performed on the KAHAN server and all plots were generated using Matplotlib.

The paper does not specify certain hyperparameters, specifically the number of epochs and batch size. This proved problematic as the interplay between these two variables affects the number of steps needed to converge. Furthermore, although five models are specified as part of the CIFAR-10 experimental set-up, it is not specified if only one model was used to generate Figure 1 in particular, or if multiple model results were averaged together.

As a result, some difference in the shape of the graph is likely due to hyperparameter tuning; specifically, we use a batch size of 125 and run on 100 epochs. Because $\frac{\# \text{ training steps}}{\text{epoch}} = \frac{\# \text{ samples}}{\text{batch size}}$, we choose the batch size and epochs such that, for 50,000 samples, we will have 25,000 training steps in total for each model run (thus spanning the width of the x-tick limits), and 400 steps per epoch (thus yielding a desirable width between data points), since data points are only logged at the end of each epoch in our Keras callback. We expect that a larger batch size should lead to a faster convergence for an otherwise fixed number of epochs. In earlier test runs (see *Diverging Models* below) a batch size of 200 was used, which may have contributed to varying convergence times; the `shuffle_pixels` function was also found to shuffle along the wrong axis, causing it to effectively randomize the location of images in the dataset rather than the data of the images themselves.

Figure 2 is produced from the source code by running `main.py` with options `-m MiniInceptionV3` and subsequently running `plot_results.py` with options `--model InceptionV3 -e 100 -b 125`.

3   Diverging Models

In the course of the assignment, we went through many iterations of models and data corruption. This allowed us to identify how different models learned, and by comparing the results against Figure 1, we looked for errors in the data corruption subroutines. Owing to convergence rates much slower than we expect from Figure 1, we discarded these models from consideration for producing our plot.
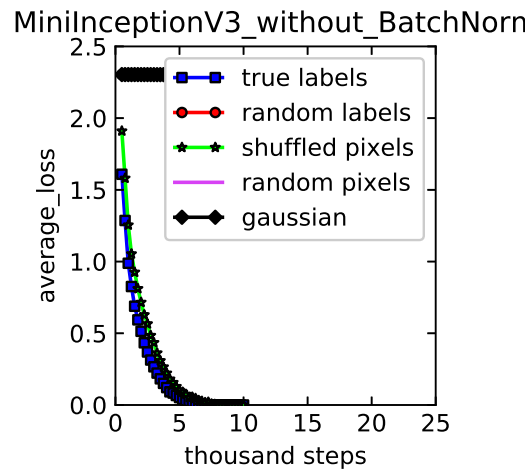
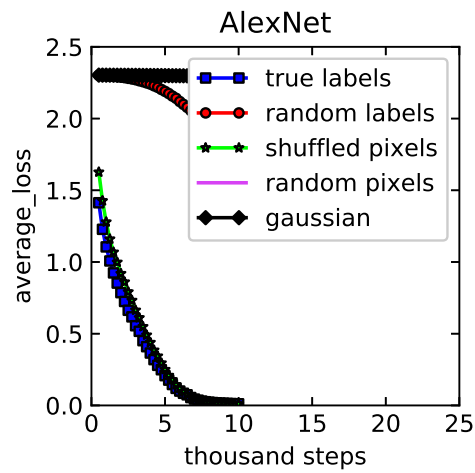Figure 3: Average losses for MiniInceptionV3 without BatchNorm.


Figure 4: Average losses for AlexNet.

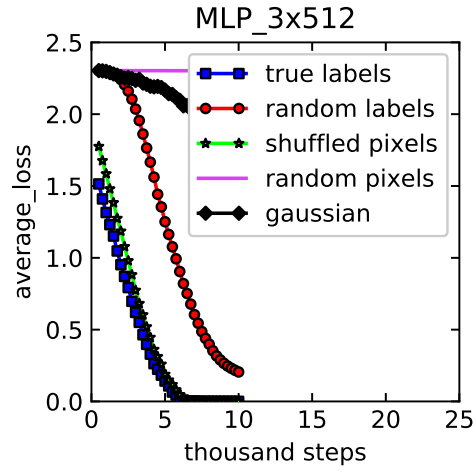Figure 5: Average losses for MLP 1x512 without BatchNorm.

Figure 6: Average losses for MLP 3x512 without BatchNorm.

Figure 7: The average losses for each data type over all of the five models: Mini-InceptionV3, MiniInceptionV3 w/o BatchNorm, AlexNet, the 1x512 MLP, and the 3x512 MLP.

4 Appendix: Code

For convenience of use, the `main.py` and `plot_results.py` scripts were augmented with command line argument parsing, giving strong user customization over which models to run on which data corruption, as well as being able to plot average losses for each data corruption type over several models.

4.1 Main Loop

main.py:

```python
import tensorflow.keras as keras
from tensorflow.keras.datasets import cifar10
from utils import *
from math import ceil
import sys
import argparse

parser = argparse.ArgumentParser()
parser.add_argument('-m', '--model_num',
    nargs = '*', default = [1, 2, 3, 4, 5])
parser.add_argument('-d', '--data_corruption_types',
    nargs = '*', default = ["true_labels", "random_labels",
                    "shuffled_pixels", "random_pixels", "gaussian"])
args = parser.parse_args()

BATCH_SIZE = 125
NUM_SAMPLES = 50000
NUM_EPOCHS = 100
STEPS_PER_EPOCH = ceil(NUM_SAMPLES / BATCH_SIZE)

def get_model(model_code):
    if model_code == 1:
        return MiniInceptionV3(
            input_shape = X.shape[1:],
            num_labels = 10
            )
    elif model_code == 2:
        return MiniInceptionV3(
```

```python
            input_shape = X.shape[1:],
            num_labels = 10,
            use_batch_norm = False
            )
    elif model_code == 3:
        return AlexNet(
            input_shape = X.shape[1:],
            num_labels = 10
        )
    elif model_code == 4:
        return MLP(
            input_shape = X.shape[1:],
            num_labels = 10,
            num_hidden_layers = 1
        )
    elif model_code == 5:
        return MLP(
            input_shape = X.shape[1:],
            num_labels = 10,
            num_hidden_layers = 3
        )

(X_train, y_train), (X_test, y_test) = cifar10.load_data()

X_train = preprocess_input(X_train)
X_test = preprocess_input(X_test)
y_train = preprocess_labels(y_train)
y_test = preprocess_labels(y_test)

inputs = {
    "true_labels": [X_train, y_train, X_test, y_test],
    "random_labels": [X_train, randomize_labels(y_train.shape[0], 10),
                X_test, randomize_labels(y_test.shape[0], 10)],
    "shuffled_pixels": [shuffle_pixels(X_train), y_train,
                shuffle_pixels(X_test), y_test],
    "random_pixels": [randomize_pixels(X_train), y_train,
                randomize_pixels(X_test), y_test],
    "gaussian": [create_gaussian_noise(X_train), y_train,
```

```
            create_gaussian_noise(X_test), y_test]}

for model_code in args.model_num:
    for corruption_type in args.data_corruption_types:
        test_X = inputs[corruption_type][2]
        test_y = inputs[corruption_type][3]
        model = get_model(model_code)
        model.compile()
        model.fit(*inputs[corruption_type],
            NUM_EPOCHS, corruption_type, BATCH_SIZE)
        model.evaluate(test_X, test_y)
```

4.2   Data Preprocessing

utils/preprocess_data.py:

```
from tensorflow.image import per_image_standardization
from tensorflow.keras.utils import to_categorical

def normalize(x_input):
    return x_input/255

def crop(x_input):
    return x_input[:, 2:-2, 2:-2, :]

def preprocess_input(x_input):
    x_out = normalize(x_input)
    x_out = crop(x_out)
    x_out = per_image_standardization(x_out)
    return x_out

def preprocess_labels(y_input):
    return to_categorical(y_input)
```

4.3   Data Corruption

utils/data_corruption.py:

```python
import numpy as np
np.random.seed(31415)

def randomize_labels(num_labels, num_unique_labels):
    random_labels = np.random.randint(low=0,
                                      high=num_unique_labels,
                                      size=num_labels)
    categorical_random_labels = np.eye(num_unique_labels)[random_labels]
    return categorical_random_labels

def shuffle_image_pixels(image_data):
    shape = image_data.shape
    shuffled_image_pixels = image_data.reshape(shape[0]*shape[1],shape[2])
    np.random.seed(31415)
    shuffled_image_pixels = np.apply_along_axis(np.random.permutation, 1, shuffl
    shuffled_image_pixels = shuffled_image_pixels.reshape(shape[0],shape[1],shap
    return shuffled_image_pixels

def shuffle_pixels(normalized_pixel_data):
    shuffled_pixel_data = np.copy(normalized_pixel_data)
    for i in range(len(shuffled_pixel_data)):
        shuffled_pixel_data[i] = shuffle_image_pixels(shuffled_pixel_data[i])
    return shuffled_pixel_data

def randomize_pixels(normalized_pixel_data):
    randomized_pixel_data = np.random.uniform(low = 0,
        high = 1.0,
        size=(normalized_pixel_data.shape))
    return randomized_pixel_data

def create_gaussian_noise(normalized_pixel_data):
    pixel_data = normalized_pixel_data.numpy().astype('float32')
    mean = pixel_data.mean(axis=0)
    std = pixel_data.std(axis=0)
    noisy_pixel_data = np.random.normal(mean, std,
        normalized_pixel_data.shape)
    np.clip(noisy_pixel_data, a_min=0, a_max = 1,
```

```
        out=noisy_pixel_data)
    return noisy_pixel_data
```

4.4 Model Definitions

utils/Models.py:

```python
import tensorflow.keras as keras
import os

RATE_DECAY_FACTOR_PER_EPOCH = 0.95
MOMENTUM_PARAMETER = 0.9

everything_in_dir = os.listdir(os.getcwd())
folders_in_dir = filter(
    lambda f: os.path.isdir(f) and "output" in f,
    everything_in_dir)
max_folder_num = 0
for folder in folders_in_dir:
    folder_num = folder[(folder.find("_") + 1):]
    max_folder_num = max(int(folder_num), max_folder_num)
OUTPUT_DIR = "output_{}".format(max_folder_num + 1)
if not os.path.exists(OUTPUT_DIR):
    os.mkdir(OUTPUT_DIR)

if keras.backend.image_data_format() == 'channels_first':
    CHANNEL_AXIS = 1
else:
    CHANNEL_AXIS = 3

if keras.backend.image_data_format() == 'channels_first':
    BATCHNORM_AXIS = 1
else:
    BATCHNORM_AXIS = 3

class MidtermModel:
    def __init__(self, weight_decay = None):
        self.model_name = "model"
```

```python
        self.model = keras.models.Model()
        self.initial_learning_rate = 0.1
        self.weight_decay = weight_decay

    def compile(self):

        sgd = keras.optimizers.SGD(
            learning_rate = self.initial_learning_rate,
            momentum = MOMENTUM_PARAMETER,
            nesterov = False
            )

        self.model.compile(
            loss="categorical_crossentropy",
            optimizer=sgd,
            metrics=["acc", "top_k_categorical_accuracy"])

        print("Finished compiling")
        self.model.summary()

    def fit(self, X_train, y_train, X_val, y_val,
        num_epoch, data_name, batch_size):
        run_name = "{}-{}".format(self.model_name, data_name)
        weights_file_name = "{}-weights.h5".format(run_name)
        weights_output_path = os.path.join(OUTPUT_DIR, weights_file_name)
        log_file_name = "{}.csv".format(run_name)
        log_output_path = os.path.join(OUTPUT_DIR, log_file_name)

        callbacks = [
                keras.callbacks.LearningRateScheduler(
                    self.learning_rate_schedule,
                    verbose = 0),
                keras.callbacks.ModelCheckpoint(weights_output_path,
                    monitor="acc",
                    save_best_only=True,
                    verbose=1),
                keras.callbacks.CSVLogger(
                    log_output_path)
```

```python
            ]

        self.model_log = self.model.fit(X_train, y_train,
            batch_size = batch_size,
            epochs = num_epoch,
            verbose = 1,
            validation_data = (X_val, y_val),
            callbacks = callbacks)

        return self.model_log

    def evaluate(self, X_test, y_test):
        self.model.evaluate(X_test, y_test, verbose = 1)

    def learning_rate_schedule(self, epoch_num):
        return self.initial_learning_rate * \
            (RATE_DECAY_FACTOR_PER_EPOCH)**(epoch_num)

class MiniInceptionV3(MidtermModel):
    def __init__(self, input_shape, num_labels=10, use_batch_norm = True):
        super(MiniInceptionV3, self).__init__()
        self.initial_learning_rate = 0.1
        self.model_name = "MiniInceptionV3"
        if not use_batch_norm:
            self.model_name += "_without_BatchNorm"
        self.use_batch_norm = use_batch_norm
        input_layer = keras.layers.Input(shape = input_shape)
        x = self.conv_module(input_layer, 96,
            kernel_size = (3,3),
            strides = (1,1))
        x = self.inception_module(x, 32, 32)
        x = self.inception_module(x, 32, 48)
        x = self.downsample_module(x, 80)
        x = self.inception_module(x, 112, 48)
        x = self.inception_module(x, 96, 64)
        x = self.inception_module(x, 80, 80)
        x = self.inception_module(x, 48, 96)
        x = self.downsample_module(x, 96)
```

```python
        x = self.inception_module(x, 176, 160)
        x = self.inception_module(x, 176, 160)
        x = keras.layers.GlobalAveragePooling2D(
            data_format = keras.backend.image_data_format())(x)
        x = keras.layers.Dense(
            num_labels, activation='softmax', name='predictions')(x)

        self.model = keras.models.Model(input_layer, x, name=self.model_name)

    def conv_module(self, input_layer, filters, kernel_size,
        padding='same', strides=(1, 1)):
        x = keras.layers.Conv2D(
            filters,
            kernel_size = kernel_size,
            strides=strides,
            padding=padding,
            use_bias=False)(input_layer)
        if self.use_batch_norm:
            x = keras.layers.BatchNormalization(
                axis=BATCHNORM_AXIS, scale=False)(x)
        x = keras.layers.Activation('relu')(x)
        return x

    def inception_module(self, input_layer, filters_1, filters_3):
        conv_module1 = self.conv_module(
            input_layer, filters = filters_1,
            kernel_size = (1, 1), strides = (1, 1))
        conv_module3 = self.conv_module(
            input_layer, filters = filters_3,
            kernel_size = (3, 3), strides = (1, 1))

        return keras.layers.concatenate(
            [conv_module1,
            conv_module3],
            axis = CHANNEL_AXIS)

    def downsample_module(self, input_layer, filters):
        max_pooling = keras.layers.MaxPooling2D(
```

```python
                pool_size = (3,3), strides = (2,2),
                padding = 'same')(input_layer)
        conv_module3 = self.conv_module(
            input_layer, filters,
            kernel_size = (3, 3), strides = (2,2))

        return keras.layers.concatenate(
            [conv_module3,
            max_pooling],
            axis = CHANNEL_AXIS)

class LocalResponseNormalization(keras.layers.Layer):
    '''
    Code sample adapted from "Deep Learning with Keras" by Gulli and Pal
    '''
    def __init__(self, n=5, alpha=0.0005, beta=0.75, k=2, **kwargs):
        self.n = n
        self.alpha = alpha
        self.beta = beta
        self.k = k
        super(LocalResponseNormalization, self).__init__(**kwargs)

    def build(self, input_shape):
        self.shape = input_shape
        super(LocalResponseNormalization, self).build(input_shape)

    def call(self, x, mask=None):
        if keras.backend.image_data_format() == 'channels_first':
            _, f, r, c = self.shape
        else:
            _, r, c, f = self.shape
        squared = keras.backend.square(x)
        pooled = keras.backend.pool2d(squared,
            (self.n,self.n), strides = (1,1),
            padding = "same", pool_mode = "avg")
        summed = keras.backend.sum(pooled, axis=CHANNEL_AXIS, keepdims = True)
        averaged = self.alpha * keras.backend.repeat_elements(
            summed, f, axis=CHANNEL_AXIS)
```

```python
        denom = keras.backend.pow(self.k + averaged, self.beta)
        return x / denom

    def get_output_shape_for(self, input_shape):
        return input_shape


class AlexNet(MidtermModel):
    def __init__(self, input_shape, num_labels=10):
        super(AlexNet, self).__init__()
        self.initial_learning_rate = 0.01
        self.model_name = "AlexNet"
        input_layer = keras.layers.Input(shape = input_shape)
        x = self.small_module(input_layer, filters = 96)
        x = self.small_module(x, filters = 256)
        x = keras.layers.Flatten()(x)
        x = keras.layers.Dense(384, activation = 'relu')(x)
        x = keras.layers.Dense(192, activation = 'relu')(x)
        x = keras.layers.Dense(
            num_labels,
            activation='softmax',
            name='predictions')(x)

        self.model = keras.models.Model(input_layer, x, name=self.model_name)

    def small_module(self, input_layer, filters = 96):
        x = keras.layers.Conv2D(
            filters,
            kernel_size = (5, 5),
            padding = 'valid')(input_layer)
        x = keras.layers.MaxPooling2D(
            pool_size = (3, 3), padding = "valid")(x)
        x = LocalResponseNormalization()(x)
        return x

class MLP(MidtermModel):
    def __init__(self, input_shape, num_labels=10,
                 num_hidden_layers = 1,
```

```python
                num_hidden_units = 512):
    super(MLP, self).__init__()
    self.model_name = "MLP_{}x{}".format(
        num_hidden_layers, num_hidden_units)
    input_layer = keras.layers.Input(shape = input_shape)
    x = keras.layers.Flatten()(input_layer)
    for hidden_layer in range(num_hidden_layers):
        x = keras.layers.Dense(num_hidden_units)(x)
        x = keras.layers.Activation("relu")(x)
    x = keras.layers.Dense(
        num_labels,
        activation = "softmax",
        name = "predictions")(x)
    self.model = keras.models.Model(
        input_layer, x, name=self.model_name)
```

4.5 Matplotlib configuation

`plot_results.py`:

```python
"""
Plot of learning curves
true labels: blue, with blue square dots
random labels - red, with red circular dots
shuffled pixels: green, with green star 5-point star dots
random pixels: purple, with no dots
gaussian: black, with black diamond dots
"""
import matplotlib.pyplot as plt
import pandas as pd
import os
import argparse
from math import ceil

label_markers = ['true', 'random_labels', 'shuffled',
                 'random_pixels', 'gaussian']
model_names = ["MiniInceptionV3",
               "MiniInceptionV3_without_BatchNorm",
```

```python
                 "AlexNet", "MLP_1x512", "MLP_3x512"]
model_correspondence = {
    "MiniInceptionV3": 1,
    "MiniInceptionV3_without_BatchNorm":2,
    "AlexNet":3,
    "MLP_1x512":4,
    "MLP_3x512":5
}

def get_max_folder_num():
    everything_in_dir = os.listdir(os.getcwd())
    folders_in_dir = filter(lambda f: os.path.isdir(f) and
                            "output" in f, everything_in_dir)
    max_folder_num = 1
    for folder in folders_in_dir:
        folder_num = folder[(folder.find("_") + 1):]
        max_folder_num = max(int(folder_num), max_folder_num)
    return max_folder_num


parser = argparse.ArgumentParser()
parser.add_argument("--model", nargs="*", default = model_names)
parser.add_argument("-e", "--num_epochs",
    nargs="?", type=int, default=100)
parser.add_argument("-b", "--batch_size",
    nargs="?", type=int, default=100)
parser.add_argument("--output_num", nargs = "?",
    default = get_max_folder_num())

parser.add_argument("-i", "--iterate", action = "store_true")
args = parser.parse_args()

OUTPUT_DIR = "output_{}".format(args.output_num)

def plot_results(steps_per_epoch, models = model_names,
                 plot_name="All_Models_Averaged"):
    def filter_dir_files(file_name):
        conditionals = [".csv" in file_name]
```

```python
    if ("MiniInceptionV3" in models and
        "MiniInceptionV3_without_BatchNorm" not in models and
        "BatchNorm" in file_name):
        return False
    else:
        conditionals.append(any([model_name in file_name
            for model_name in models]))
    return all(conditionals)

all_data = []


all_files_in_dir = list(filter(filter_dir_files,
                              os.listdir(OUTPUT_DIR)))

for label in label_markers:
    data_by_model = []
    for file_name in all_files_in_dir:
        if label in file_name:
            path = os.path.join(OUTPUT_DIR, file_name)
            data = pd.read_csv(path)
            data = data[:args.num_epochs]
            data_by_model.append(data)
    all_data.append(data_by_model)

all_losses = []
for data_corruption in all_data:
    avg_loss = pd.DataFrame({'average loss':
                            [0] * len(data_corruption[0])})
    avg_loss['epoch'] = avg_loss.index
    for model_results in data_corruption:
        avg_loss['average loss'] = avg_loss['average loss'] + \
            model_results["loss"]
    avg_loss['average loss'] = avg_loss['average loss'] / \
        len(data_corruption)
    all_losses.append(avg_loss)
for dataset in all_losses:
    dataset["thousand steps"] = \
```

```python
            ((dataset["epoch"] + 1)*steps_per_epoch)/1000
true_label_format = [{'c': "blue", 'marker': 's',
                        'edgecolors': 'black'},
                     {'c': "blue"}]

random_label_format = [{'c': "red", 'marker': 'o',
                          'edgecolors': 'black'},
                       {'c': "red"}]

shuffled_pixel_format = [{'c': '#00ff00', 'marker': '*',
                            'edgecolors': 'black'},
                         {'c': '#00ff00'}]

random_pixel_format = [{'c': None, 'marker': None,
                          'edgecolors': None},
                       {'c': "#D742F4"}]

gaussian_format = [{'c': 'black', 'marker': 'D',
                      'edgecolors': 'black'},
                   {'c': "black"}]

formats = [true_label_format, random_label_format,
           shuffled_pixel_format, random_pixel_format,
           gaussian_format]
legend_names = ['true labels', 'random labels',
                'shuffled pixels', 'random pixels', 'gaussian']

fig1 = plt.figure(figsize=(3,3))
i = 0
array_of_linmarks=[]
for dataset, data_format in zip(all_losses, formats):
    lin = None
    mark = None
    ax = plt.gca()
    z = 5 if i == 3 else 0
    if i != 3:
        mark = ax.scatter(dataset["thousand steps"].values[1:],
                    dataset["average loss"].values[1:],
```

```python
                            s = 15,
                            zorder = 10,
                            **data_format[0])
        lin, = ax.plot(dataset["thousand steps"].values[1:],
                    dataset["average loss"].values[1:],
                    zorder = z,
                    **data_format[1]
            )
        if mark:
            array_of_linmarks.append((lin, mark))
        else:
            array_of_linmarks.append((lin))
        i+=1
    ax.tick_params(axis = 'both', direction = 'in',
                    top = True, right = True)
    leg = ax.legend(array_of_linmarks, legend_names,
                    scatterpoints=2, framealpha = 1,
                    scatteryoffsets=[0.5],
                                    loc = 'upper right')
    leg.set_zorder(20)
    leg.get_frame().set_facecolor('w')
    ax.set_xticks([0, 5, 10, 15, 20, 25])
    ax.set_yticks([0, 0.5, 1, 1.5, 2.0, 2.5])

    plt.draw()

    plt.xlabel("thousand steps")
    plt.ylabel("average_loss")
    plt.title(plot_name)
    plt.xlim(0, 25)
    plt.ylim(0, 2.5)
    plt.tight_layout()
    print(os.path.join(OUTPUT_DIR, "{}.eps".format(plot_name)))
    print(os.path.join(OUTPUT_DIR, "{}.png".format(plot_name)))
    fig1.savefig(os.path.join(OUTPUT_DIR, "{}.eps".format(plot_name)))
    fig1.savefig(os.path.join(OUTPUT_DIR, "{}.png".format(plot_name)))

if __name__ == "__main__":
```

```python
if args.iterate:
    for model in args.model:
        plot_results(
            ceil(50000/args.batch_size),
            models = [model],
            plot_name = model)
else:
    plot_results(
        ceil(50000/args.batch_size),
        models = args.model)
```